

dog_app

February 19, 2019

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets: * Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.

- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays `human_files` and `dog_files`.

```
In [1]: import numpy as np
import random
from glob import glob

# load filenames for human and dog images
human_files = np.array(glob("/data/lfw/*/"))
dog_files = np.array(glob("/data/dog_images/*/"))

# print number of images in each dataset
print('There are %d total human images.' % len(human_files))
print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [2]: import cv2
import matplotlib.pyplot as plt
%matplotlib inline

# extract pre-trained face detector
face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

# load color (BGR) image
img = cv2.imread(human_files[0])

# Wrong detection examples (Human face in dog images)
# img = cv2.imread('/data/dog_images/train/103.Mastiff/Mastiff_06844.jpg')
# img = cv2.imread('/data/dog_images/train/059.Doberman_pinscher/Doberman_pinscher_04157')

# convert BGR image to grayscale
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# find faces in image
faces = face_cascade.detectMultiScale(gray)

# print number of faces detected in the image
print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [3]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0

def display_image(img_path):
    # load color (BGR) image
    img = cv2.imread(img_path)

    # convert BGR image to RGB for plotting
    cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

    # display the image, along with bounding box
    plt.imshow(cv_rgb)
    plt.show()
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: - Percentage of `human_files_short` with detected human face: **98.00%** - Percentage of `dog_files_short` with detected human face: **17.00%**

```
In [4]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

### Do NOT modify the code above this line. ###

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
## TODO
debug = False
human_count = 0
for human_file in human_files_short:
```

```

        if face_detector(human_file):
            human_count += 1
        elif debug:
            display_image(human_file)
    print("Percentage of human_files_short with detected human face: {:.2f}%".format(100*human_count/len(human_files_short)))

    human_count = 0
    for dog_file in dog_files_short:
        if face_detector(dog_file):
            human_count += 1
        if debug:
            #print('dog_file: ', dog_file)
            display_image(dog_file)
    print("Percentage of dog_files_short with detected human face: {:.2f}%".format(100*human_count/len(dog_files_short)))

```

Percentage of human_files_short with detected human face: 98.00%

Percentage of dog_files_short with detected human face: 17.00%

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on human_files_short and dog_files_short.

```

In [5]: ### (Optional)
        ### TODO: Test performance of another face detection algorithm.
        ### Feel free to use as many code cells as needed.

```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```

In [6]: import torch
        import torchvision.models as models

        # define VGG16 model
        VGG16 = models.vgg16(pretrained=True)

        # check if CUDA is available
        use_cuda = torch.cuda.is_available()

```

```
# move model to GPU if CUDA is available
if use_cuda:
    VGG16 = VGG16.cuda()
```

Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth
 100%|| 553433881/553433881 [00:04<00:00, 116403068.07it/s]

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [7]: from PIL import Image, ImageFile
        ImageFile.LOAD_TRUNCATED_IMAGES = True
        import torchvision.transforms as transforms

        def VGG16_predict(img_path):
            """
            Use pre-trained VGG-16 model to obtain index corresponding to
            predicted ImageNet class for image at specified path

            Args:
                img_path: path to an image

            Returns:
                Index corresponding to VGG-16 model's prediction
            """

            ## TODO: Complete the function.
            ## Load and pre-process an image from the given img_path
            ## Return the *index* of the predicted class for that image

            # Load Image
            img_bgr = Image.open(img_path)

            # Define Image Transform to make it compatible with VGG-16
            data_transform = transforms.Compose([ transforms.Resize([224, 224]),
                                                  transforms.ToTensor(),
                                                  transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                                           std=[0.229, 0.224, 0.225])])
```

```

# Transform Image
# unsqueeze - Feed input as a batch of size 1 (1 image in the batch)
# img_tensor = data_transform(img_bgr).unsqueeze(0)
# Discard the transparent, alpha channel (that's the :3) and add the batch dimension
img_tensor = data_transform(img_bgr)[:3,:,:].unsqueeze(0)

# Move input to GPU if available
if use_cuda:
    img_tensor = img_tensor.cuda()

# Feed Forward the Image via VGG-16
output = VGG16(img_tensor)

# Get class with max probability
_, pred = torch.max(output, 1)

#if debug:
    #print(f"Prediction for {img_path} is: ", pred)
    #plt.imshow(img_bgr)

return pred # predicted class index

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns True if a dog is detected in an image (and False if not).

```

In [8]: ### returns "True" if a dog is detected in the image stored at img_path
# Return true/false
def dog_detector(img_path):
    ## TODO: Complete the function.
    predicted_class = VGG16_predict(img_path)
    if predicted_class >= 151 and predicted_class <= 268:
        return True
    return False

test_file = dog_files[random.randint(0,99)]
print(f"dog_detector({test_file}):", dog_detector(test_file))

```

```
dog_detector(/data/dog_images/train/103.Mastiff/Mastiff_06828.jpg): True
```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?

- What percentage of the images in `dog_files_short` have a detected dog?

Answer: - Percentage of `human_files_short` with detected dog: **1.00%** - Percentage of `dog_files_short` with detected dog: **100.00%**

```
In [9]: ### TODO: Test the performance of the dog_detector function  
### on the images in human_files_short and dog_files_short.
```

```
dog_count = 0  
for human_file in human_files_short:  
    if dog_detector(human_file):  
        dog_count += 1  
        if debug:  
            display_image(human_file)  
print("Percentage of human_files_short with detected dog: {:.2f}%".format(100*dog_count/len(human_files_short)))  
  
dog_count = 0  
for dog_file in dog_files_short:  
    if dog_detector(dog_file):  
        dog_count += 1  
    else:  
        if debug:  
            display_image(dog_file)  
print("Percentage of dog_files_short with detected dog: {:.2f}%".format(100*dog_count/len(dog_files_short)))
```

Percentage of `human_files_short` with detected dog: 1.00%

Percentage of `dog_files_short` with detected dog: 100.00%

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [10]: ### (Optional)  
### TODO: Report the performance of another pre-trained network.  
### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [11]: import os
import numpy as np
import torchvision.transforms as transforms
import torch

from torchvision import datasets
from PIL import Image
from torch.utils.data.sampler import SubsetRandomSampler

### TODO: Write data loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes
data_dir = '/data/dog_images/'
```

```

train_dir = os.path.join(data_dir, 'train/')
valid_dir = os.path.join(data_dir, 'valid/')
test_dir = os.path.join(data_dir, 'test/')

# Get the dog class labels
classes = os.listdir(train_dir)

# Transform - Random Resized Crop
data_transform1 = transforms.Compose([transforms.RandomResizedCrop(224),
                                     transforms.RandomHorizontalFlip(),
                                     transforms.ToTensor(),
                                     transforms.Normalize(mean=[0.485, 0.456, 0.406],st
data_transform2 = transforms.Compose([transforms.Resize([224, 224]),
                                     transforms.ToTensor(),
                                     transforms.Normalize(mean=[0.485, 0.456, 0.406],st

# get data sets
train_data = datasets.ImageFolder(train_dir, transform=data_transform1)
valid_data = datasets.ImageFolder(valid_dir, transform=data_transform2)
test_data = datasets.ImageFolder(test_dir, transform=data_transform2)

# print out some data stats
num_train = len(train_data)
num_valid = len(valid_data)
num_test = len(test_data)
num_classes = len(classes)
print('Num training images: ', num_train)
print('Num validation images: ', num_valid)
print('Num test images: ', num_test)
print('Num classes: ', num_classes)

# define dataloader parameters
batch_size = 20
num_workers = 0

# setup the loaders
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, num_worke
valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size, num_worke
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size, num_workers
loaders_scratch = {'train': train_loader,
                  'valid': valid_loader,
                  'test': test_loader}

```

```

Num training images: 6680
Num validation images: 835
Num test images: 836
Num classes: 133

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer: - Resizing - **Training:** I used RandomResizedCrop(224) based on what is used in several other pretrained torchvision models. I also normalized the images using the mean and standard deviation used by the pretrained models. - **Validation/Testing:** I used transforms.Resize([224, 224]) so as to lose as little information (vs RandomCrop) as possible during evaluation. - Tensor Size - I used a **224x224** image size again based on the what was used by torchvision models like VGG16. The size was not too small and allowed to me to still get RandomCrops of the original image with enough variety to help with dataset augmentation. - Dataset Augmentation - **Yes.** - I used RandomResizedCrop and RandomHorizontalFlip to augment the data (training on more variations of the dataset) and also to avoid overfitting. - In addition, I tried using RandomRotation(10) on the training set datatransform. But, the model accuracy decreased! Most likely because I did not train the model for a long time (>30 epochs).

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```
In [12]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        self.conv1 = nn.Conv2d(3, 32, 3, stride=2, padding=1)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=2, padding=1)
        self.conv3 = nn.Conv2d(64, 128, 3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(7 * 7 * 128, 512)
        self.fc2 = nn.Linear(512, num_classes)
        self.dropout = nn.Dropout(p=0.2)

    def forward(self, x):
        ## Define forward behavior

        # convolution layers
        x = self.pool(F.relu((self.conv1(x))))
        x = self.pool(F.relu((self.conv2(x))))
        x = self.pool(F.relu((self.conv3(x))))
        #print('Shape: ', x.shape)
```

```

        # flatten
        x = x.view(-1, 7 * 7 * 128)

        # pass to fully connected layers
        x = self.dropout(x)
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = self.fc2(x)
        return x

    """### You so NOT have to modify the code below this line. ###"""

    # instantiate the CNN
    model_scratch = Net()
    print(model_scratch)

    # move tensors to GPU if CUDA is available
    if use_cuda:
        model_scratch.cuda()

Net(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=6272, out_features=512, bias=True)
  (fc2): Linear(in_features=512, out_features=133, bias=True)
  (dropout): Dropout(p=0.2)
)

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer: - Input Image Size: - 224x224 - Convolution Layers: - Wanted to reduce the (x, y) dimensions to (7x7). Inline with well established models. - Depth: Tried 64 and later increased it to 128 to get over 10% accuracy. - Layer 1: Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1)) - MaxPool: MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False) - Layer 2: Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1)) - MaxPool: MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False) - Layer 3: Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)) - Modified stride here so as to achieve the desired (x, y) dimensions of (7, 7) after all the convolution layers. Reducing (x,y) to (3, 3) did not train well. - MaxPool: MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False) - Fully Connected Layers: - Used 2 fully connected layers with ReLu activation and a dropout of 20% - Layer 1: (fc1): Linear(in_features=6272, out_features=512, bias=True) - Layer 2: (fc2): Linear(in_features=512, out_features=133, bias=True) - Loss: CrossEntropyLoss - Optimizer: SGD (with learning rate = 0.05). Tried Adam too which worked equally well. - Output Size (num classes): - 133 - **Test Accuracy (on 1 run):** 20% (171/836)

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
In [13]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.05)
```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```
In [14]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""
        # initialize tracker for minimum validation loss
        valid_loss_min = np.Inf
        print_every = 100

        for epoch in range(1, n_epochs+1):
            # initialize variables to monitor training and validation loss
            train_loss = 0.0
            valid_loss = 0.0

            #####
            # train the model #
            #####
            model.train()
            for batch_idx, (data, target) in enumerate(loaders['train']):
                # move to GPU
                if use_cuda:
                    data, target = data.cuda(), target.cuda()
                ## find the loss and update the model parameters accordingly
                ## record the average training loss, using something like
                ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

                # Reset gradients
                optimizer.zero_grad()

                # Feed forward and compute loss
                output = model(data)

                # compute loss
                loss = criterion(output, target)
```

```

        # optimizer step
        loss.backward()
        optimizer.step()

    train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))

    if debug and batch_idx % print_every == 0:
        print(f'\tEpoch #{epoch}, Iteration #{batch_idx+1}, Loss: {train_loss}')

    #####
    # validate the model #
    #####
    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['valid']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## update the average validation loss

        # Feed forward and compute loss
        output = model(data)
        loss = criterion(output, target)
        valid_loss += ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
    ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss < valid_loss_min:
        print('\tValidation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.
            format(
                valid_loss_min,
                valid_loss))
        valid_loss_min = valid_loss
        torch.save(model.state_dict(), save_path)

    print("\n")

    # return trained model
    return model

# train the model
num_epochs = 30
debug=True

```

```

model_file = 'models/model_scratch.pt'

# load the previously saved model if available, to get a warm start
if os.path.isfile(model_file):
    model_scratch.load_state_dict(torch.load(model_file))

model_scratch = train(num_epochs, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, model_file)

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load(model_file))

print('Done Training!')

Epoch #1, Iteration #1, Loss: 3.222787857055664
Epoch #1, Iteration #101, Loss: 3.415236473083496
Epoch #1, Iteration #201, Loss: 3.4306204319000244
Epoch #1, Iteration #301, Loss: 3.431774377822876
Epoch: 1      Training Loss: 3.416890      Validation Loss: 3.859243
Validation loss decreased (inf --> 3.859243). Saving model ...

Epoch #2, Iteration #1, Loss: 2.9762368202209473
Epoch #2, Iteration #101, Loss: 3.3477654457092285
Epoch #2, Iteration #201, Loss: 3.3717944622039795
Epoch #2, Iteration #301, Loss: 3.3874828815460205
Epoch: 2      Training Loss: 3.391319      Validation Loss: 3.578955
Validation loss decreased (3.859243 --> 3.578955). Saving model ...

Epoch #3, Iteration #1, Loss: 2.5569567680358887
Epoch #3, Iteration #101, Loss: 3.248321533203125
Epoch #3, Iteration #201, Loss: 3.2913012504577637
Epoch #3, Iteration #301, Loss: 3.321340322494507
Epoch: 3      Training Loss: 3.346957      Validation Loss: 3.736090

Epoch #4, Iteration #1, Loss: 3.172574520111084
Epoch #4, Iteration #101, Loss: 3.3154990673065186
Epoch #4, Iteration #201, Loss: 3.2750895023345947
Epoch #4, Iteration #301, Loss: 3.28889536857605
Epoch: 4      Training Loss: 3.301481      Validation Loss: 3.657772

Epoch #5, Iteration #1, Loss: 3.534672260284424
Epoch #5, Iteration #101, Loss: 3.2341620922088623
Epoch #5, Iteration #201, Loss: 3.243776321411133

```

Epoch #5, Iteration #301, Loss: 3.2569363117218018
Epoch: 5 Training Loss: 3.258054 Validation Loss: 3.548971
Validation loss decreased (3.578955 --> 3.548971). Saving model ...

Epoch #6, Iteration #1, Loss: 2.587562084197998
Epoch #6, Iteration #101, Loss: 3.173037052154541
Epoch #6, Iteration #201, Loss: 3.1951847076416016
Epoch #6, Iteration #301, Loss: 3.2215211391448975
Epoch: 6 Training Loss: 3.220563 Validation Loss: 3.570187

Epoch #7, Iteration #1, Loss: 3.163062572479248
Epoch #7, Iteration #101, Loss: 3.1206462383270264
Epoch #7, Iteration #201, Loss: 3.1501362323760986
Epoch #7, Iteration #301, Loss: 3.1777756214141846
Epoch: 7 Training Loss: 3.177391 Validation Loss: 3.596332

Epoch #8, Iteration #1, Loss: 2.9891936779022217
Epoch #8, Iteration #101, Loss: 3.1256346702575684
Epoch #8, Iteration #201, Loss: 3.15085506439209
Epoch #8, Iteration #301, Loss: 3.1663148403167725
Epoch: 8 Training Loss: 3.169110 Validation Loss: 3.546840
Validation loss decreased (3.548971 --> 3.546840). Saving model ...

Epoch #9, Iteration #1, Loss: 3.3368678092956543
Epoch #9, Iteration #101, Loss: 3.061596632003784
Epoch #9, Iteration #201, Loss: 3.0721092224121094
Epoch #9, Iteration #301, Loss: 3.0885891914367676
Epoch: 9 Training Loss: 3.096459 Validation Loss: 3.571249

Epoch #10, Iteration #1, Loss: 3.0263917446136475
Epoch #10, Iteration #101, Loss: 3.0734403133392334
Epoch #10, Iteration #201, Loss: 3.0812151432037354
Epoch #10, Iteration #301, Loss: 3.092982053756714
Epoch: 10 Training Loss: 3.097695 Validation Loss: 3.786886

Epoch #11, Iteration #1, Loss: 2.676487445831299
Epoch #11, Iteration #101, Loss: 3.042048454284668
Epoch #11, Iteration #201, Loss: 3.037799835205078
Epoch #11, Iteration #301, Loss: 3.04154634475708
Epoch: 11 Training Loss: 3.040271 Validation Loss: 3.558982

Epoch #12, Iteration #1, Loss: 2.902878522872925
Epoch #12, Iteration #101, Loss: 2.993830442428589
Epoch #12, Iteration #201, Loss: 3.0025572776794434
Epoch #12, Iteration #301, Loss: 3.0164196491241455
Epoch: 12 Training Loss: 3.025769 Validation Loss: 3.443315
Validation loss decreased (3.546840 --> 3.443315). Saving model ...

Epoch #13, Iteration #1, Loss: 2.888326406478882
Epoch #13, Iteration #101, Loss: 2.921394109725952
Epoch #13, Iteration #201, Loss: 2.944779634475708
Epoch #13, Iteration #301, Loss: 2.953549385070801
Epoch: 13 Training Loss: 2.958822 Validation Loss: 3.531041

Epoch #14, Iteration #1, Loss: 3.128679037094116
Epoch #14, Iteration #101, Loss: 2.943382501602173
Epoch #14, Iteration #201, Loss: 2.9561848640441895
Epoch #14, Iteration #301, Loss: 2.965667486190796
Epoch: 14 Training Loss: 2.973637 Validation Loss: 3.507233

Epoch #15, Iteration #1, Loss: 2.4618849754333496
Epoch #15, Iteration #101, Loss: 2.951514959335327
Epoch #15, Iteration #201, Loss: 2.9548890590667725
Epoch #15, Iteration #301, Loss: 2.9454963207244873
Epoch: 15 Training Loss: 2.946701 Validation Loss: 3.532644

Epoch #16, Iteration #1, Loss: 2.6506834030151367
Epoch #16, Iteration #101, Loss: 2.831879138946533
Epoch #16, Iteration #201, Loss: 2.8525216579437256
Epoch #16, Iteration #301, Loss: 2.8761136531829834
Epoch: 16 Training Loss: 2.890655 Validation Loss: 3.808220

Epoch #17, Iteration #1, Loss: 3.4228038787841797
Epoch #17, Iteration #101, Loss: 2.8707993030548096
Epoch #17, Iteration #201, Loss: 2.8619353771209717
Epoch #17, Iteration #301, Loss: 2.887681007385254
Epoch: 17 Training Loss: 2.891411 Validation Loss: 3.538228

Epoch #18, Iteration #1, Loss: 2.1164374351501465
Epoch #18, Iteration #101, Loss: 2.794858694076538
Epoch #18, Iteration #201, Loss: 2.8487281799316406
Epoch #18, Iteration #301, Loss: 2.8851191997528076
Epoch: 18 Training Loss: 2.894778 Validation Loss: 3.632461

Epoch #19, Iteration #1, Loss: 3.13734769821167
Epoch #19, Iteration #101, Loss: 2.733853578567505
Epoch #19, Iteration #201, Loss: 2.790557622909546
Epoch #19, Iteration #301, Loss: 2.803269147872925
Epoch: 19 Training Loss: 2.814364 Validation Loss: 3.555188

Epoch #20, Iteration #1, Loss: 2.2411577701568604
Epoch #20, Iteration #101, Loss: 2.853034496307373
Epoch #20, Iteration #201, Loss: 2.8291540145874023
Epoch #20, Iteration #301, Loss: 2.823007822036743
Epoch: 20 Training Loss: 2.828866 Validation Loss: 3.516930

Epoch #21, Iteration #1, Loss: 3.198730945587158
Epoch #21, Iteration #101, Loss: 2.759795904159546
Epoch #21, Iteration #201, Loss: 2.756319046020508
Epoch #21, Iteration #301, Loss: 2.7767271995544434
Epoch: 21 Training Loss: 2.794479 Validation Loss: 3.516197

Epoch #22, Iteration #1, Loss: 2.129220962524414
Epoch #22, Iteration #101, Loss: 2.702435255050659
Epoch #22, Iteration #201, Loss: 2.74430775642395
Epoch #22, Iteration #301, Loss: 2.77054500579834
Epoch: 22 Training Loss: 2.781784 Validation Loss: 3.528630

Epoch #23, Iteration #1, Loss: 2.7940332889556885
Epoch #23, Iteration #101, Loss: 2.6588544845581055
Epoch #23, Iteration #201, Loss: 2.6730949878692627
Epoch #23, Iteration #301, Loss: 2.701401710510254
Epoch: 23 Training Loss: 2.716979 Validation Loss: 3.605093

Epoch #24, Iteration #1, Loss: 2.7360739707946777
Epoch #24, Iteration #101, Loss: 2.608131170272827
Epoch #24, Iteration #201, Loss: 2.6754415035247803
Epoch #24, Iteration #301, Loss: 2.7130370140075684
Epoch: 24 Training Loss: 2.712008 Validation Loss: 3.712774

Epoch #25, Iteration #1, Loss: 3.4913086891174316
Epoch #25, Iteration #101, Loss: 2.632678985595703
Epoch #25, Iteration #201, Loss: 2.6637189388275146
Epoch #25, Iteration #301, Loss: 2.691129446029663

Epoch: 25 Training Loss: 2.695329 Validation Loss: 3.613812

Epoch #26, Iteration #1, Loss: 3.2120487689971924
Epoch #26, Iteration #101, Loss: 2.65832781791687
Epoch #26, Iteration #201, Loss: 2.655308961868286
Epoch #26, Iteration #301, Loss: 2.677370309829712

Epoch: 26 Training Loss: 2.684457 Validation Loss: 3.471386

Epoch #27, Iteration #1, Loss: 1.8021968603134155
Epoch #27, Iteration #101, Loss: 2.688344717025757
Epoch #27, Iteration #201, Loss: 2.675527811050415
Epoch #27, Iteration #301, Loss: 2.6795949935913086

Epoch: 27 Training Loss: 2.672729 Validation Loss: 3.577964

Epoch #28, Iteration #1, Loss: 2.474621295928955
Epoch #28, Iteration #101, Loss: 2.513354778289795
Epoch #28, Iteration #201, Loss: 2.560866117477417
Epoch #28, Iteration #301, Loss: 2.5960423946380615

Epoch: 28 Training Loss: 2.616052 Validation Loss: 3.627453

Epoch #29, Iteration #1, Loss: 1.9054960012435913
Epoch #29, Iteration #101, Loss: 2.5672268867492676
Epoch #29, Iteration #201, Loss: 2.590578556060791
Epoch #29, Iteration #301, Loss: 2.5955138206481934

Epoch: 29 Training Loss: 2.606864 Validation Loss: 3.620353

Epoch #30, Iteration #1, Loss: 2.1952805519104004
Epoch #30, Iteration #101, Loss: 2.505970001220703
Epoch #30, Iteration #201, Loss: 2.549293279647827
Epoch #30, Iteration #301, Loss: 2.5775768756866455

Epoch: 30 Training Loss: 2.594254 Validation Loss: 3.672189

Done Training!

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [15]: def test(loaders, model, criterion, use_cuda):
```

```

# monitor test loss and accuracy
test_loss = 0.
correct = 0.
total = 0.

model.eval()
for batch_idx, (data, target) in enumerate(loaders['test']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the loss
    loss = criterion(output, target)
    # update average test loss
    test_loss += ((1 / (batch_idx + 1)) * (loss.data - test_loss))
    # convert output probabilities to predicted class
    pred = output.data.max(1, keepdim=True)[1]
    # compare predictions to true label
    correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
    total += data.size(0)

print('Test Loss: {:.6f}\n'.format(test_loss))

print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
    100. * correct / total, correct, total))

```

In [16]: *# call test function*

```
test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

Test Loss: 3.390768

Test Accuracy: 20% (171/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

In [17]: import os
import numpy as np
import torchvision.transforms as transforms
import torch

from torchvision import datasets
from PIL import Image, ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True
from torch.utils.data.sampler import SubsetRandomSampler

### TODO: Write data loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes
data_dir = '/data/dog_images/'
train_dir = os.path.join(data_dir, 'train/')
valid_dir = os.path.join(data_dir, 'valid/')
test_dir = os.path.join(data_dir, 'test/')

# Get the dog class labels
classes = os.listdir(train_dir)

# Transform - Random Resized Crop
data_transform_train = transforms.Compose([transforms.RandomResizedCrop(299),
                                           transforms.RandomHorizontalFlip(),
                                           transforms.ToTensor(),
                                           transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])]),
data_transform_test = transforms.Compose([transforms.Resize([299, 299]),
                                           transforms.ToTensor(),
                                           transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])]),

# get data sets
train_data = datasets.ImageFolder(train_dir, transform=data_transform_train)
valid_data = datasets.ImageFolder(valid_dir, transform=data_transform_test)
test_data = datasets.ImageFolder(test_dir, transform=data_transform_test)

# print out some data stats
num_train = len(train_data)
num_valid = len(valid_data)
num_test = len(test_data)
num_classes = len(classes)
print('Num training images: ', num_train)
print('Num validation images: ', num_valid)
print('Num test images: ', num_test)
print('Num classes: ', num_classes)

# define dataloader parameters
batch_size = 20
num_workers = 0

```

```

# setup the loaders
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, num_workers=
valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size, num_workers=
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size, num_workers=
loaders_transfer = {'train': train_loader,
                    'valid': valid_loader,
                    'test': test_loader}

```

```

Num training images: 6680
Num validation images: 835
Num test images: 836
Num classes: 133

```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [18]: import torch
import torchvision.models as models
import torch.nn as nn

import torch.optim as optim

# TODO: Specify model architecture
model_transfer = models.inception_v3(pretrained=True)

# check if CUDA is available
use_cuda = torch.cuda.is_available()

# Freeze training for all "features" layers
for param in model_transfer.parameters():
    param.requires_grad = False

# Replace fully connected layer
model_transfer.fc = nn.Linear(2048, num_classes)

# Replace Auxillary Classifier layer
model_transfer.AuxLogits.fc = nn.Linear(768, num_classes)

# Print modified model
print(model_transfer)

if use_cuda:
    model_transfer = model_transfer.cuda()

```

Downloading: "https://download.pytorch.org/models/inception_v3_google-1a9a5a14.pth" to /root/.torch
100%|| 108857766/108857766 [00:01<00:00, 55897450.75it/s]

```
Inception3(
  (Conv2d_1a_3x3): BasicConv2d(
    (conv): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), bias=False)
    (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (Conv2d_2a_3x3): BasicConv2d(
    (conv): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (Conv2d_2b_3x3): BasicConv2d(
    (conv): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (Conv2d_3b_1x1): BasicConv2d(
    (conv): Conv2d(64, 80, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(80, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (Conv2d_4a_3x3): BasicConv2d(
    (conv): Conv2d(80, 192, kernel_size=(3, 3), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (Mixed_5b): InceptionA(
    (branch1x1): BasicConv2d(
      (conv): Conv2d(192, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch5x5_1): BasicConv2d(
      (conv): Conv2d(192, 48, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn): BatchNorm2d(48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch5x5_2): BasicConv2d(
      (conv): Conv2d(48, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2), bias=False)
      (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_1): BasicConv2d(
      (conv): Conv2d(192, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_2): BasicConv2d(
      (conv): Conv2d(64, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_3): BasicConv2d(
      (conv): Conv2d(96, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

```

```

        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(192, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(Mixed_5c): InceptionA(
    (branch1x1): BasicConv2d(
        (conv): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch5x5_1): BasicConv2d(
        (conv): Conv2d(256, 48, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch5x5_2): BasicConv2d(
        (conv): Conv2d(48, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_1): BasicConv2d(
        (conv): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_2): BasicConv2d(
        (conv): Conv2d(64, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_3): BasicConv2d(
        (conv): Conv2d(96, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(Mixed_5d): InceptionA(
    (branch1x1): BasicConv2d(
        (conv): Conv2d(288, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch5x5_1): BasicConv2d(
        (conv): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(48, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch5x5_2): BasicConv2d(
        (conv): Conv2d(48, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2), bias=False)

```



```

        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_1): BasicConv2d(
        (conv): Conv2d(288, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_2): BasicConv2d(
        (conv): Conv2d(64, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_3): BasicConv2d(
        (conv): Conv2d(96, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(288, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(Mixed_6a): InceptionB(
    (branch3x3): BasicConv2d(
        (conv): Conv2d(288, 384, kernel_size=(3, 3), stride=(2, 2), bias=False)
        (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_1): BasicConv2d(
        (conv): Conv2d(288, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_2): BasicConv2d(
        (conv): Conv2d(64, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch3x3dbl_3): BasicConv2d(
        (conv): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(Mixed_6b): InceptionC(
    (branch1x1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7_1): BasicConv2d(
        (conv): Conv2d(768, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7_2): BasicConv2d(
        (conv): Conv2d(128, 128, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)

```

```

    (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7_3): BasicConv2d(
  (conv): Conv2d(128, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
  (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_1): BasicConv2d(
  (conv): Conv2d(768, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_2): BasicConv2d(
  (conv): Conv2d(128, 128, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
  (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_3): BasicConv2d(
  (conv): Conv2d(128, 128, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
  (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_4): BasicConv2d(
  (conv): Conv2d(128, 128, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
  (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_5): BasicConv2d(
  (conv): Conv2d(128, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
  (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch_pool): BasicConv2d(
  (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
)
(Mixed_6c): InceptionC(
  (branch1x1): BasicConv2d(
    (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7_1): BasicConv2d(
    (conv): Conv2d(768, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7_2): BasicConv2d(
    (conv): Conv2d(160, 160, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7_3): BasicConv2d(
    (conv): Conv2d(160, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
)

```

```

(branch7x7dbl_1): BasicConv2d(
  (conv): Conv2d(768, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_2): BasicConv2d(
  (conv): Conv2d(160, 160, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
  (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_3): BasicConv2d(
  (conv): Conv2d(160, 160, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
  (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_4): BasicConv2d(
  (conv): Conv2d(160, 160, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
  (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_5): BasicConv2d(
  (conv): Conv2d(160, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
  (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch_pool): BasicConv2d(
  (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
)
(Mixed_6d): InceptionC(
  (branch1x1): BasicConv2d(
    (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7_1): BasicConv2d(
    (conv): Conv2d(768, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7_2): BasicConv2d(
    (conv): Conv2d(160, 160, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7_3): BasicConv2d(
    (conv): Conv2d(160, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7dbl_1): BasicConv2d(
    (conv): Conv2d(768, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7dbl_2): BasicConv2d(
    (conv): Conv2d(160, 160, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)

```

```

        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7dbl_3): BasicConv2d(
        (conv): Conv2d(160, 160, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7dbl_4): BasicConv2d(
        (conv): Conv2d(160, 160, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7dbl_5): BasicConv2d(
        (conv): Conv2d(160, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(Mixed_6e): InceptionC(
    (branch1x1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7_1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7_2): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7_3): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7dbl_1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7dbl_2): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
    (branch7x7dbl_3): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
    )
)

```

```

(branch7x7dbl_4): BasicConv2d(
  (conv): Conv2d(192, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
  (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch7x7dbl_5): BasicConv2d(
  (conv): Conv2d(192, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
  (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch_pool): BasicConv2d(
  (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
)
(AuxLogits): InceptionAux(
  (conv0): BasicConv2d(
    (conv): Conv2d(768, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (conv1): BasicConv2d(
    (conv): Conv2d(128, 768, kernel_size=(5, 5), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(768, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (fc): Linear(in_features=768, out_features=133, bias=True)
)
(Mixed_7a): InceptionD(
  (branch3x3_1): BasicConv2d(
    (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch3x3_2): BasicConv2d(
    (conv): Conv2d(192, 320, kernel_size=(3, 3), stride=(2, 2), bias=False)
    (bn): BatchNorm2d(320, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7x3_1): BasicConv2d(
    (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7x3_2): BasicConv2d(
    (conv): Conv2d(192, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0, 3), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7x3_3): BasicConv2d(
    (conv): Conv2d(192, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3, 0), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch7x7x3_4): BasicConv2d(
    (conv): Conv2d(192, 192, kernel_size=(3, 3), stride=(2, 2), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
)

```

```

    )
)
(Mixed_7b): InceptionE(
  (branch1x1): BasicConv2d(
    (conv): Conv2d(1280, 320, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(320, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch3x3_1): BasicConv2d(
    (conv): Conv2d(1280, 384, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch3x3_2a): BasicConv2d(
    (conv): Conv2d(384, 384, kernel_size=(1, 3), stride=(1, 1), padding=(0, 1), bias=False)
    (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch3x3_2b): BasicConv2d(
    (conv): Conv2d(384, 384, kernel_size=(3, 1), stride=(1, 1), padding=(1, 0), bias=False)
    (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch3x3dbl_1): BasicConv2d(
    (conv): Conv2d(1280, 448, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(448, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch3x3dbl_2): BasicConv2d(
    (conv): Conv2d(448, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch3x3dbl_3a): BasicConv2d(
    (conv): Conv2d(384, 384, kernel_size=(1, 3), stride=(1, 1), padding=(0, 1), bias=False)
    (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch3x3dbl_3b): BasicConv2d(
    (conv): Conv2d(384, 384, kernel_size=(3, 1), stride=(1, 1), padding=(1, 0), bias=False)
    (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch_pool): BasicConv2d(
    (conv): Conv2d(1280, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(Mixed_7c): InceptionE(
  (branch1x1): BasicConv2d(
    (conv): Conv2d(2048, 320, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(320, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )
  (branch3x3_1): BasicConv2d(
    (conv): Conv2d(2048, 384, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
  )

```

```

)
(branch3x3_2a): BasicConv2d(
  (conv): Conv2d(384, 384, kernel_size=(1, 3), stride=(1, 1), padding=(0, 1), bias=False)
  (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch3x3_2b): BasicConv2d(
  (conv): Conv2d(384, 384, kernel_size=(3, 1), stride=(1, 1), padding=(1, 0), bias=False)
  (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch3x3dbl_1): BasicConv2d(
  (conv): Conv2d(2048, 448, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn): BatchNorm2d(448, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch3x3dbl_2): BasicConv2d(
  (conv): Conv2d(448, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch3x3dbl_3a): BasicConv2d(
  (conv): Conv2d(384, 384, kernel_size=(1, 3), stride=(1, 1), padding=(0, 1), bias=False)
  (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch3x3dbl_3b): BasicConv2d(
  (conv): Conv2d(384, 384, kernel_size=(3, 1), stride=(1, 1), padding=(1, 0), bias=False)
  (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
(branch_pool): BasicConv2d(
  (conv): Conv2d(2048, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True, track_running_stats=True)
)
)
(fc): Linear(in_features=2048, out_features=133, bias=True)
)

```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: I wanted to try some new architectures and chose Inception_v3 as it is known to perform well, heard it mentioned several times but actually not seen it in action. I soon realized that this model had an **auxiliary fully connected layer** that was used only in training. I replaced both the output and auxiliary fully connected layers with my own that matched with the number of output classes for our classification problem. I froze all model parameter other than the ones relating to the 2 fully connected layers that I added.

- Tensor Size
- I had to use 299x299 image size based on the input requirements for Inception_v3 model.
- Training: I used RandomResizedCrop(229). I also normalized the images using the mean and standard deviation used by the pretrained models.
- Validation/Testing: I used transforms.Resize([299, 299]) so as to lose as little information (vs RandomCrop) as possible during evaluation.

- Dataset Augmentation
 - Yes.
 - I used RandomResizedCrop and RandomHorizontalFlip to augment the data (training on more variations of the dataset) and also to avoid overfitting.
- Test Accuracy (on 1 run): 78% (659/836)

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [19]: criterion_transfer = nn.CrossEntropyLoss()
         params_to_update = []
         for param in model_transfer.parameters():
             if param.requires_grad:
                 params_to_update.append(param)
         optimizer_transfer = optim.Adam(params_to_update, lr=0.01)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [20]: import os.path

         def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path, is_inception):
             """returns trained model"""
             # initialize tracker for minimum validation loss
             valid_loss_min = np.Inf
             print_every = 100

             for epoch in range(1, n_epochs+1):
                 # initialize variables to monitor training and validation loss
                 train_loss = 0.0
                 valid_loss = 0.0

                 #####
                 # train the model #
                 #####
                 model.train()
                 for batch_idx, (data, target) in enumerate(loaders['train']):
                     # move to GPU
                     if use_cuda:
                         data, target = data.cuda(), target.cuda()
                     ## find the loss and update the model parameters accordingly
                     ## record the average training loss, using something like
                     ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
```



```

# Reset gradients
optimizer.zero_grad()

# Feed forward and compute loss
if is_inception:
    # handle auxillary output from inception
    # https://pytorch.org/tutorials/beginner/finetuning_torchvision_models_
    output, aux_output = model(data)
    loss1 = criterion(output, target)
    loss2 = criterion(aux_output, target)
    loss = loss1 + 0.4*loss2
else:
    output = model(data)
    loss = criterion(output, target)

# optimizer step
loss.backward()
optimizer.step()

train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))

if debug and batch_idx % print_every == 0:
    print(f'\tEpoch #{epoch}, Iteration #{batch_idx+1}, Loss: {train_loss}'

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## update the average validation loss

    # Feed forward and compute loss
    # Even in inception, there is no auxillary output during model evaluation
    # https://github.com/pytorch/vision/blob/master/torchvision/models/inception
    output = model(data)
    loss = criterion(output, target)

    valid_loss += ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss

```

```

    ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss < valid_loss_min:
        print('\tValidation loss decreased ({:.6f} --> {:.6f}). Saving model ...'
              valid_loss_min,
              valid_loss))
        valid_loss_min = valid_loss
        torch.save(model.state_dict(), save_path)

    print("\n")

    # return trained model
    return model

In [21]: # train the model
        debug=True
        num_epochs = 30
        model_file = 'models/model_transfer.pt'

        # load the previously saved model if available, to get a warm start
        if os.path.isfile(model_file):
            model_transfer.load_state_dict(torch.load(model_file))

        model_transfer = train(num_epochs, loaders_transfer, model_transfer, optimizer_transfer,
                               criterion_transfer, use_cuda, model_file, True)

        # load the model that got the best validation accuracy (uncomment the line below)
        model_transfer.load_state_dict(torch.load(model_file))

        print('Done Training!')

Epoch #1, Iteration #1, Loss: 2.5857138633728027
Epoch #1, Iteration #101, Loss: 6.0715532302856445
Epoch #1, Iteration #201, Loss: 6.115310192108154
Epoch #1, Iteration #301, Loss: 6.2081217765808105
Epoch: 1      Training Loss: 6.268450      Validation Loss: 2.021797
Validation loss decreased (inf --> 2.021797). Saving model ...

Epoch #2, Iteration #1, Loss: 7.233248710632324
Epoch #2, Iteration #101, Loss: 5.946088790893555
Epoch #2, Iteration #201, Loss: 6.140623569488525
Epoch #2, Iteration #301, Loss: 6.265612602233887
Epoch: 2      Training Loss: 6.317700      Validation Loss: 2.139870

Epoch #3, Iteration #1, Loss: 6.745226860046387

```

Epoch #3, Iteration #101, Loss: 6.348630428314209
Epoch #3, Iteration #201, Loss: 6.434607982635498
Epoch #3, Iteration #301, Loss: 6.634353160858154
Epoch: 3 Training Loss: 6.648694 Validation Loss: 2.014069
Validation loss decreased (2.021797 --> 2.014069). Saving model ...

Epoch #4, Iteration #1, Loss: 4.540031433105469
Epoch #4, Iteration #101, Loss: 6.114386081695557
Epoch #4, Iteration #201, Loss: 6.2734503746032715
Epoch #4, Iteration #301, Loss: 6.387295722961426
Epoch: 4 Training Loss: 6.471759 Validation Loss: 1.817088
Validation loss decreased (2.014069 --> 1.817088). Saving model ...

Epoch #5, Iteration #1, Loss: 9.699653625488281
Epoch #5, Iteration #101, Loss: 6.729641437530518
Epoch #5, Iteration #201, Loss: 6.790037155151367
Epoch #5, Iteration #301, Loss: 6.856613636016846
Epoch: 5 Training Loss: 6.972218 Validation Loss: 1.965765

Epoch #6, Iteration #1, Loss: 9.631362915039062
Epoch #6, Iteration #101, Loss: 6.873438835144043
Epoch #6, Iteration #201, Loss: 7.3469438552856445
Epoch #6, Iteration #301, Loss: 7.179884433746338
Epoch: 6 Training Loss: 7.249672 Validation Loss: 1.930143

Epoch #7, Iteration #1, Loss: 4.078765392303467
Epoch #7, Iteration #101, Loss: 7.176174163818359
Epoch #7, Iteration #201, Loss: 7.156558036804199
Epoch #7, Iteration #301, Loss: 7.09710168838501
Epoch: 7 Training Loss: 7.160083 Validation Loss: 2.118811

Epoch #8, Iteration #1, Loss: 6.949888229370117
Epoch #8, Iteration #101, Loss: 6.753422737121582
Epoch #8, Iteration #201, Loss: 6.999551773071289
Epoch #8, Iteration #301, Loss: 7.178799629211426
Epoch: 8 Training Loss: 7.180485 Validation Loss: 2.223028

Epoch #9, Iteration #1, Loss: 13.026325225830078
Epoch #9, Iteration #101, Loss: 6.869229316711426
Epoch #9, Iteration #201, Loss: 7.133728981018066
Epoch #9, Iteration #301, Loss: 7.373859405517578
Epoch: 9 Training Loss: 7.469930 Validation Loss: 2.613167

Epoch #10, Iteration #1, Loss: 6.99279260635376
Epoch #10, Iteration #101, Loss: 7.021251201629639
Epoch #10, Iteration #201, Loss: 7.2922821044921875
Epoch #10, Iteration #301, Loss: 7.384606838226318
Epoch: 10 Training Loss: 7.462847 Validation Loss: 2.484450

Epoch #11, Iteration #1, Loss: 6.730443000793457
Epoch #11, Iteration #101, Loss: 7.525222301483154
Epoch #11, Iteration #201, Loss: 7.333653926849365
Epoch #11, Iteration #301, Loss: 7.407731056213379
Epoch: 11 Training Loss: 7.322731 Validation Loss: 2.490700

Epoch #12, Iteration #1, Loss: 9.245223045349121
Epoch #12, Iteration #101, Loss: 7.923200607299805
Epoch #12, Iteration #201, Loss: 7.800239086151123
Epoch #12, Iteration #301, Loss: 7.8751091957092285
Epoch: 12 Training Loss: 7.812053 Validation Loss: 2.266948

Epoch #13, Iteration #1, Loss: 3.8581244945526123
Epoch #13, Iteration #101, Loss: 7.609640598297119
Epoch #13, Iteration #201, Loss: 7.722017288208008
Epoch #13, Iteration #301, Loss: 7.710282325744629
Epoch: 13 Training Loss: 7.670198 Validation Loss: 2.213556

Epoch #14, Iteration #1, Loss: 8.212874412536621
Epoch #14, Iteration #101, Loss: 7.797274589538574
Epoch #14, Iteration #201, Loss: 7.799099445343018
Epoch #14, Iteration #301, Loss: 8.031720161437988
Epoch: 14 Training Loss: 8.115748 Validation Loss: 2.375376

Epoch #15, Iteration #1, Loss: 3.9355766773223877
Epoch #15, Iteration #101, Loss: 7.430359363555908
Epoch #15, Iteration #201, Loss: 7.507734298706055
Epoch #15, Iteration #301, Loss: 7.606751441955566
Epoch: 15 Training Loss: 7.670272 Validation Loss: 2.337484

Epoch #16, Iteration #1, Loss: 10.928091049194336
Epoch #16, Iteration #101, Loss: 7.7007365226745605
Epoch #16, Iteration #201, Loss: 7.755185127258301
Epoch #16, Iteration #301, Loss: 7.874589443206787

Epoch: 16 Training Loss: 7.879434 Validation Loss: 2.383183

Epoch #17, Iteration #1, Loss: 11.433257102966309
Epoch #17, Iteration #101, Loss: 8.118145942687988
Epoch #17, Iteration #201, Loss: 7.930103302001953
Epoch #17, Iteration #301, Loss: 8.17491340637207

Epoch: 17 Training Loss: 8.165871 Validation Loss: 2.447206

Epoch #18, Iteration #1, Loss: 9.608397483825684
Epoch #18, Iteration #101, Loss: 7.685166835784912
Epoch #18, Iteration #201, Loss: 7.895205974578857
Epoch #18, Iteration #301, Loss: 8.220281600952148

Epoch: 18 Training Loss: 8.183845 Validation Loss: 2.500071

Epoch #19, Iteration #1, Loss: 5.463583469390869
Epoch #19, Iteration #101, Loss: 7.917751312255859
Epoch #19, Iteration #201, Loss: 8.120160102844238
Epoch #19, Iteration #301, Loss: 8.120307922363281

Epoch: 19 Training Loss: 8.105686 Validation Loss: 2.584613

Epoch #20, Iteration #1, Loss: 8.257987022399902
Epoch #20, Iteration #101, Loss: 7.950231075286865
Epoch #20, Iteration #201, Loss: 8.159234046936035
Epoch #20, Iteration #301, Loss: 8.262646675109863

Epoch: 20 Training Loss: 8.242275 Validation Loss: 2.808084

Epoch #21, Iteration #1, Loss: 8.799139022827148
Epoch #21, Iteration #101, Loss: 7.865368366241455
Epoch #21, Iteration #201, Loss: 7.935146331787109
Epoch #21, Iteration #301, Loss: 8.023098945617676

Epoch: 21 Training Loss: 8.080071 Validation Loss: 2.623898

Epoch #22, Iteration #1, Loss: 7.115983009338379
Epoch #22, Iteration #101, Loss: 7.763071060180664
Epoch #22, Iteration #201, Loss: 7.8278093338012695
Epoch #22, Iteration #301, Loss: 8.110808372497559

Epoch: 22 Training Loss: 8.100648 Validation Loss: 2.758134

Epoch #23, Iteration #1, Loss: 8.458528518676758
Epoch #23, Iteration #101, Loss: 8.57793140411377
Epoch #23, Iteration #201, Loss: 8.696394920349121

Epoch #23, Iteration #301, Loss: 8.624100685119629
Epoch: 23 Training Loss: 8.616224 Validation Loss: 2.863217

Epoch #24, Iteration #1, Loss: 2.8071908950805664
Epoch #24, Iteration #101, Loss: 8.155203819274902
Epoch #24, Iteration #201, Loss: 8.369111061096191
Epoch #24, Iteration #301, Loss: 8.24611759185791
Epoch: 24 Training Loss: 8.352648 Validation Loss: 3.169943

Epoch #25, Iteration #1, Loss: 6.092153549194336
Epoch #25, Iteration #101, Loss: 8.009735107421875
Epoch #25, Iteration #201, Loss: 8.258747100830078
Epoch #25, Iteration #301, Loss: 8.280364036560059
Epoch: 25 Training Loss: 8.277253 Validation Loss: 3.206868

Epoch #26, Iteration #1, Loss: 4.326980113983154
Epoch #26, Iteration #101, Loss: 8.205707550048828
Epoch #26, Iteration #201, Loss: 7.744072437286377
Epoch #26, Iteration #301, Loss: 7.997096061706543
Epoch: 26 Training Loss: 8.133477 Validation Loss: 2.451951

Epoch #27, Iteration #1, Loss: 6.755497455596924
Epoch #27, Iteration #101, Loss: 8.681264877319336
Epoch #27, Iteration #201, Loss: 8.749443054199219
Epoch #27, Iteration #301, Loss: 8.931979179382324
Epoch: 27 Training Loss: 8.934000 Validation Loss: 2.904748

Epoch #28, Iteration #1, Loss: 8.173095703125
Epoch #28, Iteration #101, Loss: 8.541572570800781
Epoch #28, Iteration #201, Loss: 8.58267879486084
Epoch #28, Iteration #301, Loss: 8.554693222045898
Epoch: 28 Training Loss: 8.633287 Validation Loss: 2.707103

Epoch #29, Iteration #1, Loss: 7.8459625244140625
Epoch #29, Iteration #101, Loss: 8.552529335021973
Epoch #29, Iteration #201, Loss: 8.319947242736816
Epoch #29, Iteration #301, Loss: 8.400060653686523
Epoch: 29 Training Loss: 8.385900 Validation Loss: 2.662087

Epoch #30, Iteration #1, Loss: 8.343340873718262
Epoch #30, Iteration #101, Loss: 8.973581314086914

```
Epoch #30, Iteration #201, Loss: 9.019980430603027
Epoch #30, Iteration #301, Loss: 8.91418170928955
Epoch: 30      Training Loss: 8.946994      Validation Loss: 2.459717
```

Done Training!

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [22]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 2.109863
```

```
Test Accuracy: 78% (659/836)
```

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [30]: def load_input_image_as_tensor(img_path):
        image = Image.open(img_path).convert('RGB')

        # discard the transparent, alpha channel (that's the :3) and add the batch dimension
        image_tensor = data_transform_test(image)[:3,:,:,].unsqueeze(0)

        return image_tensor

In [38]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        ### TODO: Write a function that takes a path to an image as input
        ### and returns the dog breed that is predicted by the model.

        # list of class names by index, i.e. a name can be accessed like class_names[0]
        class_names = [item[4:].replace("_", " ") for item in loaders_transfer['train'].dataset]

        def predict_breed_transfer(img_path):
            # load the image and return the predicted breed
            image_tensor = load_input_image_as_tensor(img_path)
            if use_cuda:
                image_tensor = image_tensor.cuda()
```

```

output = model_transfer(image_tensor)

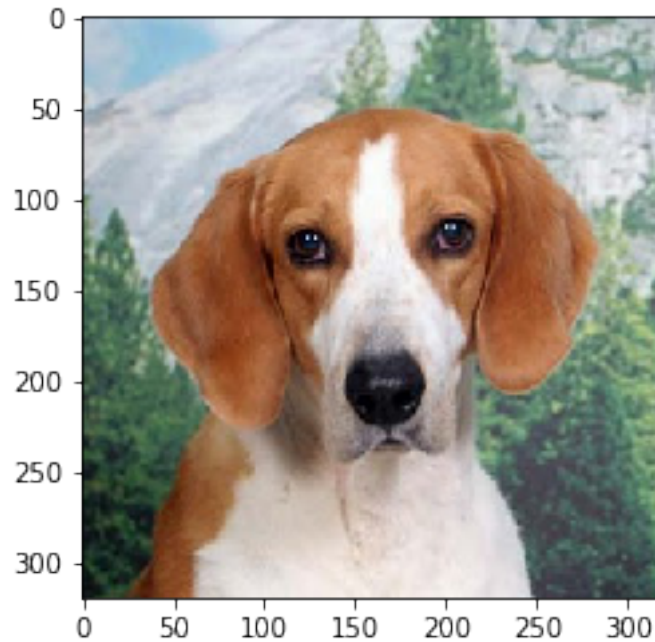
# display the image
img = cv2.imread(img_path)
# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
plt.imshow(cv_rgb)
plt.show()

# class index
pred = output.data.max(1, keepdim=True)[1]

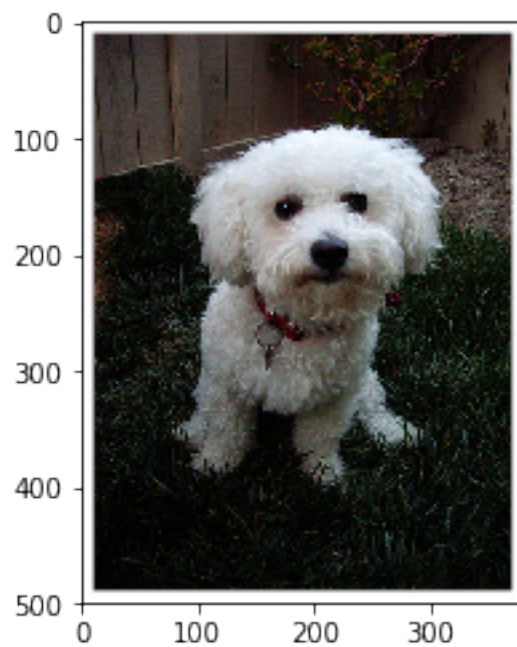
return class_names[pred]

num_samples = 5
for i in range(0, num_samples):
    test_file = dog_files[random.randint(0, len(dog_files))]
    print(f"^^^predict_breed_transfer({test_file}):", predict_breed_transfer(test_file))


```




```
^^^predict_breed_transfer(/data/dog_images/train/007.American_foxhound/American_foxhound_00474.j
```

^^^predict_breed_transfer(/data/dog_images/train/024.Bichon_frise/Bichon_frise_01698.jpg): Bichon



```
^^predict_breed_transfer(/data/dog_images/train/090.Italian_greyhound/Italian_greyhound_06132.j
```



```
^^predict_breed_transfer(/data/dog_images/train/082.Havanese/Havanese_05610.jpg): Havanese
```





Sample Human Output

```
^^predict_breed_transfer(/data/dog_images/train/021.Belgian_sheepdog/Belgian_sheepdog_01506.jpg
```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm

In [39]: *### TODO: Write your algorithm.*

Feel free to use as many code cells as needed.

```
def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    dog_breed = predict_breed_transfer(img_path)
    if dog_detector(img_path):
        print(f'Dog in "{img_path}" is a "{dog_breed}"')
    elif face_detector(img_path):
        print(f'Human in "{img_path}" resembles the dog breed "{dog_breed}"')
    else:
        print(f'Oops! Neither dog nor human is detected in file "{img_path}"')
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

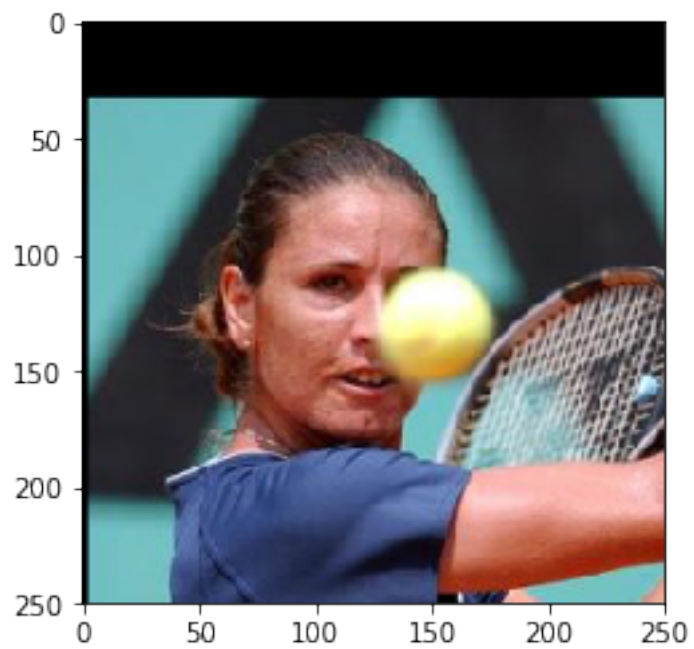
Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement) - Output Expectation - On the task of dog vs human face detection - model performed well and in line with my expectation (based on evaluations in the beginning of this notebook). - On the task of dog breed classification, the model seems to be not doing great (only ~70% accuracy) and performed worse than my expectation.

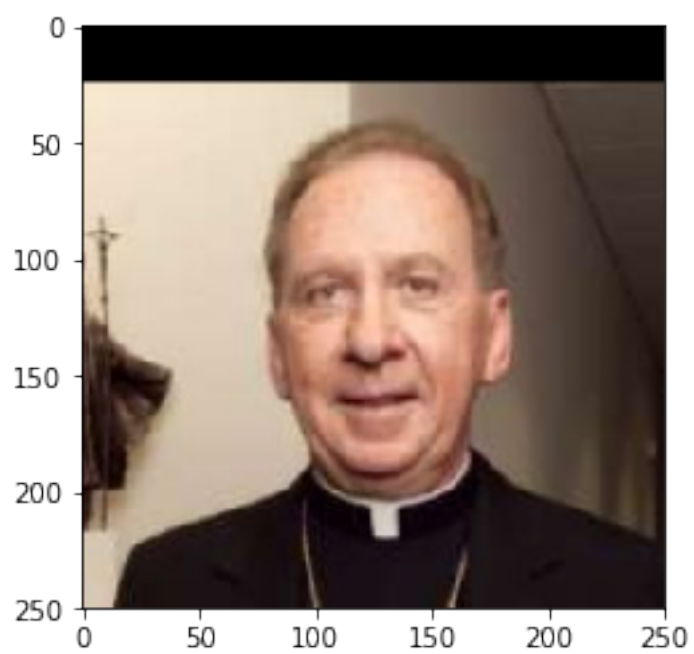
- Areas for improvement - **Train on bigger dataset.** This dataset has only 6680/133 \approx 50 samples per dog breed and some of these breeds are so close to each other that even humans will have difficulty with the task of breed identification. - **Data augmentation.** Try more data augmentation techniques (eg: rotation) and train for longer duration for convergence. - **Tuning Hyperparameters.** Learning rate, optimizer, dropout, weight initialization, normalization etc - **Ensemble Models.** Combine the results from multiple classification models (as a voting scheme) and see if this approach does better.

```
In [40]: ## TODO: Execute your algorithm from Step 6 on
         ## at least 6 images on your computer.
         ## Feel free to use as many code cells as needed.


         ## suggested code, below
num_samples = 3
for i in range(0, num_samples):
    test_file = human_files[random.randint(0, len(human_files))]
    run_app(test_file)
for i in range(0, num_samples):
    test_file = dog_files[random.randint(0, len(dog_files))]
    run_app(test_file)
```



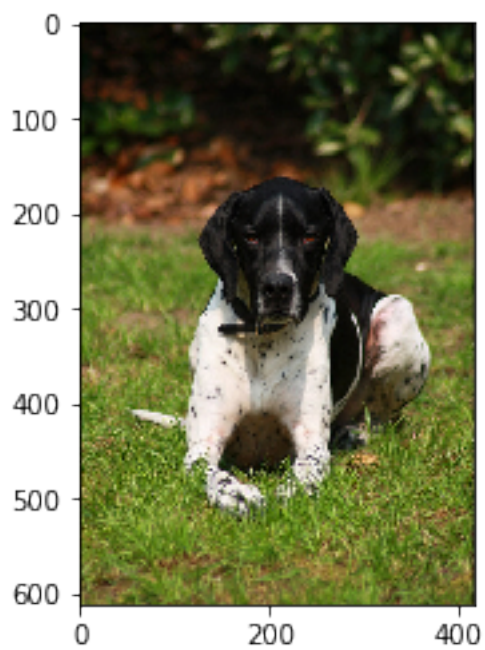
Human in "/data/lfw/Rita_Grande/Rita_Grande_0002.jpg" resembles the dog breed "Boykin spaniel"



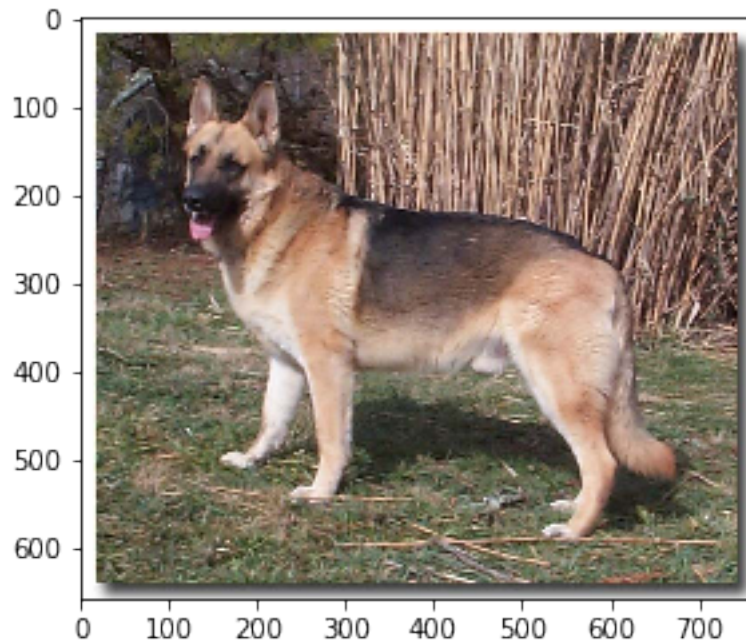
Human in `"/data/lfw/Thomas_OBrien/Thomas_OBrien_0003.jpg"` resembles the dog breed `"English springer"`




Human in `"/data/lfw/James_Williams/James_Williams_0001.jpg"` resembles the dog breed `"Afghan hound"`



Dog in `"/data/dog_images/train/122.Pointer/Pointer_07835.jpg"` is a `"Pointer"`

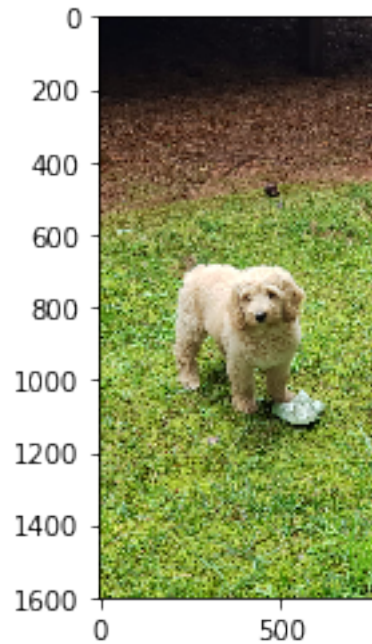


Dog in `"/data/dog_images/train/071.German_shepherd_dog/German_shepherd_dog_04944.jpg"` is a `"German_shepherd_dog"`




Dog in "/data/dog_images/train/050.Chinese_shar-pei/Chinese_shar-pei_03542.jpg" is a "Chinese sh

```
In [41]: for img_file in os.listdir('images'):
          img_path = os.path.join('images', img_file)
          run_app(img_path)
```




Dog in "images/73100552-FC55-4669-80B3-B6482F6E0CA8.JPG" is a "Bichon frise"



Human in "images/511F2544-99AD-48FF-88D3-482E59F6DCF9.jpg" resembles the dog breed "Kerry blue t



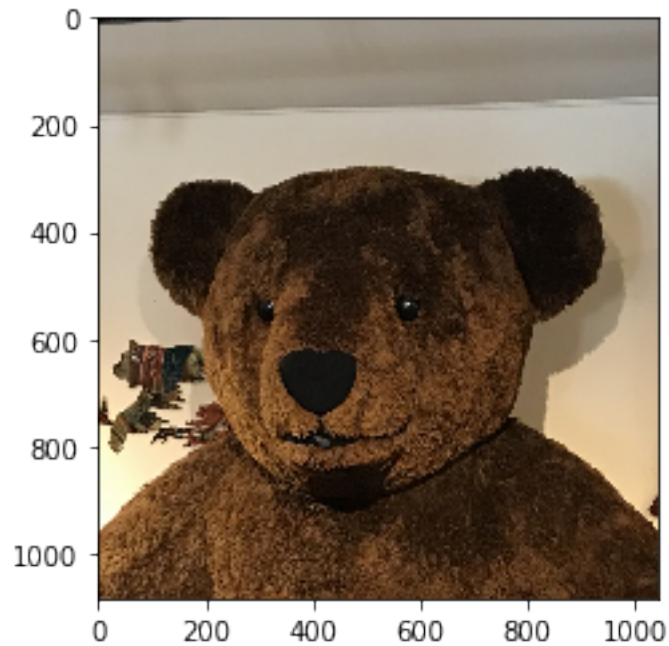
Dog in "images/Brittany_02625.jpg" is a "Brittany"




Dog in "images/Welsh_springer_spaniel_08203.jpg" is a "Welsh springer spaniel"




Dog in "images/Labrador_retriever_06449.jpg" is a "Labrador retriever"



Oops! Neither dog nor human is detected in file "images/F10C0AA7-8859-452D-B31C-E413C67AA03C.jpg"



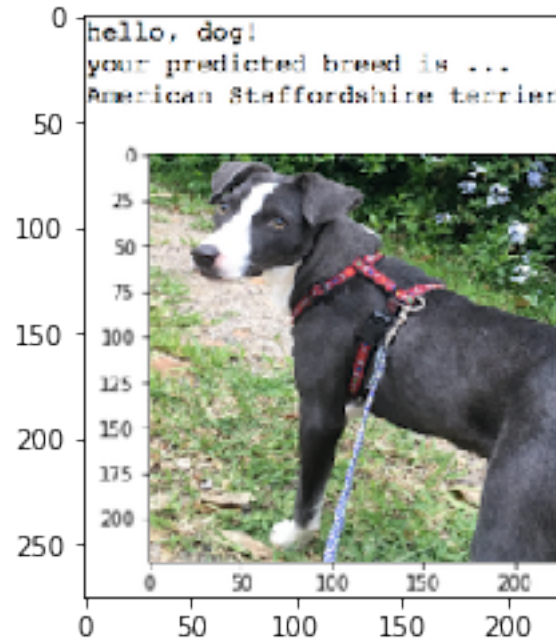
Human in "images/sample_human_output.png" resembles the dog breed "Basenji"




Dog in "images/Labrador_retriever_06457.jpg" is a "Labrador retriever"




Dog in "images/American_water_spaniel_00648.jpg" is a "Boykin spaniel"




Dog in "images/sample_dog_output.png" is a "Smooth fox terrier"



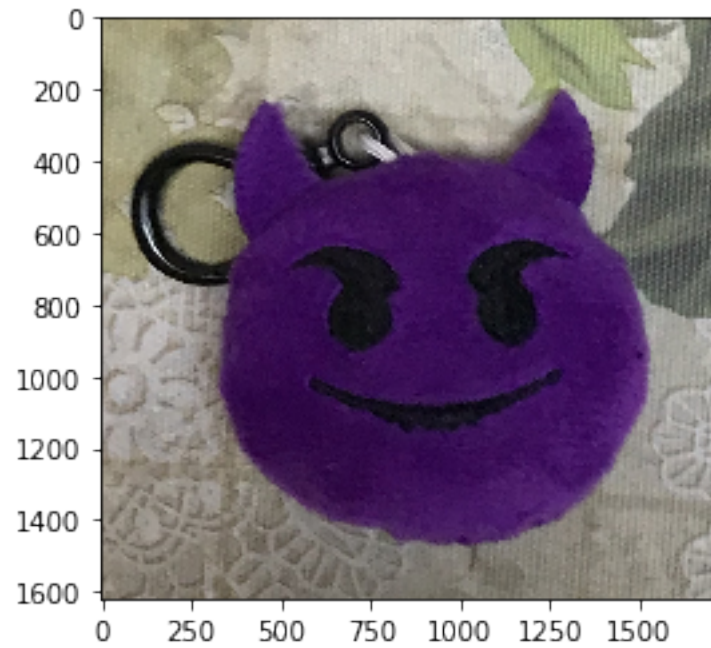
Oops! Neither dog nor human is detected in file "images/sample_cnn.png"



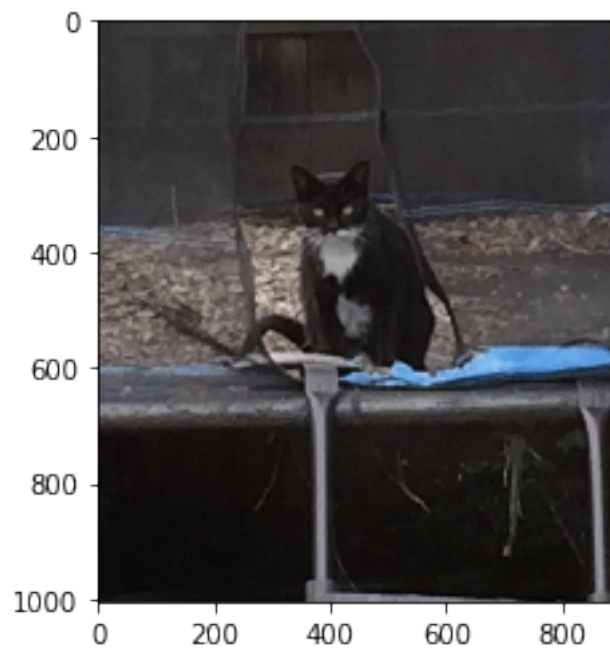
Dog in "images/Curly-coated_retriever_03896.jpg" is a "Curly-coated retriever"



Dog in "images/Labrador_retriever_06455.jpg" is a "Labrador retriever"



Oops! Neither dog nor human is detected in file "images/CC0E0128-09CF-412A-AB4F-912C7E847880.JPE"



```
Dog in "images/D782B1C8-7A59-4B16-B8A4-CE4D99433C50.jpg" is a "Border collie"
```

```
In [ ]:
```